

Vatsal Sharan

CONTACT INFORMATION	SAL 220, 941 Bloom Walk, Los Angeles, CA 90089	e-mail: vsharan@usc.edu
APPOINTMENTS	University of Southern California <i>Assistant Professor,</i> <i>Department of Computer Science</i>	Fall 2021 – present
	Massachusetts Institute of Technology <i>Norbert Wiener Postdoctoral Associate,</i> <i>Institute for Data, Systems & Society</i>	2020 – 2021
EDUCATION	Stanford University <i>Ph.D. in Electrical Engineering</i> <i>Advisor: Gregory Valiant, Dept. of Computer Science</i>	2014 – 2020
	Indian Institute of Technology Kanpur <i>B.Tech. in Electrical Engineering</i>	2010 – 2014
DISTINCTIONS	<ul style="list-style-type: none">• Amazon Research Award• NSF CAREER Award• Amazon Research Award• Best Paper Award at 35th Conference on Learning Theory (COLT)• Norbert Wiener Postdoctoral Fellowship, MIT• Outstanding reviewer at ICML'19, NeurIPS'21• Invited to China Theory Week, Tsinghua University• Director's Gold Medal for best all-round performance and leadership in graduating class, IIT Kanpur• Best Final Year Project in Electrical Engineering, IIT Kanpur• Honda Young Engineer and Scientist Award	<div>2024</div> <div>2023</div> <div>2022</div> <div>2022</div> <div>2020</div> <div></div> <div>2018</div> <div>2014</div> <div>2014</div> <div>2013</div>
SELECTED PUBLICATIONS	(most papers in theory venues have alphabetical author ordering)	
	<ol style="list-style-type: none">When is Multicalibration Post-Processing Necessary? Dutch Hansen, Siddhartha Devic, Preetum Nakkiran, Vatsal Sharan <i>Neural Information Processing Systems (NeurIPS) 2024</i>Transformers Learn Higher-Order Optimization Methods for In-Context Learning: A Study with Linear Models Deqing Fu, Tian-Qi Chen, Robin Jia, Vatsal Sharan <i>Neural Information Processing Systems (NeurIPS) 2024</i> <i>Best Paper Award at SoCal NLP Symposium 2023</i>Pre-trained Large Language Models Use Fourier Features to Compute Addition Tianyi Zhou, Deqing Fu, Vatsal Sharan, Robin Jia <i>Neural Information Processing Systems (NeurIPS) 2024</i>Optimal Multiclass U-Calibration Error and Beyond Haipeng Luo, Spandan Senapati, Vatsal Sharan <i>Neural Information Processing Systems (NeurIPS) 2024</i>Transductive Sample Complexities Are Compact Julian Asilis, Siddhartha Devic, Shaddin Dughmi, Vatsal Sharan, Shang-Hua Teng <i>Neural Information Processing Systems (NeurIPS) 2024</i>	

6. **On the Statistical Complexity of Sample Amplification**
 Brian Axelrod, Shivam Garg, Yanjun Han, Vatsal Sharan, Gregory Valiant
Annals of Statistics, 2024
7. **Regularization and Optimal Multiclass Learning**
 Julian Asilis, Siddhartha Devic, Shaddin Dughmi, Vatsal Sharan, Shang-Hua Teng
Conference on Learning Theory (COLT) 2024
8. **Stability and Group Fairness in Ranking with Uncertain Predictions**
 Siddhartha Devic, Aleksandra Korolova, David Kempe, Vatsal Sharan
International Conference on Machine Learning (ICML) 2024
Non-archival at Symposium on Foundations of Responsible Computing (FORC 2024)
9. **Mitigating Simplicity Bias in Deep Learning for Improved OOD Generalization and Robustness**
 Bhavya Vasudeva, Kameron Shahabi, Vatsal Sharan
Transactions on Machine Learning Research (TMLR) 2024
10. **Fairness in Matching under Uncertainty**
 Siddhartha Devic, David Kempe, Vatsal Sharan, Aleksandra Korolova
International Conference on Machine Learning (ICML) 2023
ACM Conference on Equity & Access in Algorithms, Mechanisms, & Optimization (EAAMO'23)
11. **NeuroSketch: A Neural Network Method for Fast and Approximate Evaluation of Range Aggregate Queries**
 Sepanta Zeighami, Vatsal Sharan, Cyrus Shahabi
ACM Special Interest Group on Management of Data Conference (SIGMOD) 2023
12. **Efficient Convex Optimization Requires Superlinear Memory**
 Annie Marsden, Vatsal Sharan, Aaron Sidford, Gregory Valiant
Conference on Learning Theory (COLT), 2022 (Best Paper Award)
Invited to IJCAI 2023 Sister Conference Notable Paper Track
Journal of the ACM (JACM), 2024
13. **Efficient Gradient Methods for Objectives with Multiple Scales**
 Jon Kelner, Annie Marsden, Vatsal Sharan, Aaron Sidford, Gregory Valiant, Honglin Yuan
Conference on Learning Theory (COLT) 2022
14. **Multicalibrated Partitions for Importance Weights**
 Parikshit Gopalan, Omer Reingold, Vatsal Sharan, Udi Wieder
Algorithmic Learning Theory (ALT) 2022
15. **Omnipredictors**
 Parikshit Gopalan, Adam Tauman Kalai, Omer Reingold, Vatsal Sharan, Udi Wieder
Innovations in Theoretical Computer Science (ITCS) 2022
16. **Modular versus Monolithic Task Formulations in Neural Networks Learning**
 Atish Agarwala, Abhimanyu Das, Brendan Juba, Rina Panigrahy, Vatsal Sharan, Xin Wang, Qiuyi Zhang
International Conference on Learning Representations (ICLR) 2021
17. **Sample Amplification: Increasing Dataset Size even when Learning is Impossible**
 Brian Axelrod, Shivam Garg, Vatsal Sharan, Gregory Valiant
International Conference on Machine Learning (ICML) 2020
18. **PIDForest: Anomaly detection via Partial Identification**
 Parikshit Gopalan, Vatsal Sharan, Udi Wieder
Neural Information Processing Systems (NeurIPS) 2019 (Spotlight presentation)
19. **Fast and Accurate Low-Rank Factorization of Compressively-Sensed Data**
 Vatsal Sharan, Kai Sheng Tai, Peter Bailis, Gregory Valiant
International Conference on Machine Learning (ICML) 2019

20. **Memory-sample Tradeoffs for Linear Regression with Small Error**
Vatsal Sharan, Aaron Sidford, Gregory Valiant
Symposium on the Theory of Computing (STOC) 2019
21. **Recovery Guarantees for Quadratic Tensors with Limited Observations**
Hongyang Zhang, Vatsal Sharan, Moses Charikar and Yingyu Liang
Artificial Intelligence & Statistics (AISTATS) 2019
22. **A Spectral View of Adversarially Robust Features**
Shivam Garg, Vatsal Sharan, Brian Zhang, Gregory Valiant
Neural Information Processing Systems (NeurIPS) 2018 (Spotlight presentation)
23. **Efficient Anomaly Detection via Matrix Sketching**
Vatsal Sharan, Parikshit Gopalan, Udi Wieder
Neural Information Processing Systems (NeurIPS) 2018
24. **Prediction with a Short Memory**
Vatsal Sharan, Sham Kakade, Percy Liang, Gregory Valiant
Symposium on the Theory of Computing (STOC) 2018
25. **Sketching Linear Classifiers over Data Streams**
Kai Sheng Tai, Vatsal Sharan, Peter Bailis, Gregory Valiant
ACM SIGMOD Conference on Management of Data (SIGMOD) 2018
26. **Moment-Based Quantile Sketches for Efficient High Cardinality Aggregation Queries**
Edward Gan, Jialin Ding, Kai Sheng Tai, Vatsal Sharan, Peter Bailis
Conference on Very Large Data Bases (VLDB) 2018
27. **Learning Overcomplete HMMs**
Vatsal Sharan, Sham Kakade, Percy Liang, Gregory Valiant
Neural Information Processing Systems (NeurIPS) 2017
28. **Orthogonalized Alternating Least Squares: A Theoretically Principled Tensor Factorization Algorithm for Practical Use**
Vatsal Sharan, Gregory Valiant
International Conference on Machine Learning (ICML) 2017

STUDENTS

PhD students

Bhavya Vasudeva (started Fall'21)
 Siddhartha Devic (started Fall'21, co-advised with Aleksandra Korolova)
 Julian Asilis (started Fall'22)
 Deqing Fu (started Fall'22, co-advised with Robin Jia)
 Devansh Gupta (started Fall'23, co-advised with Meisam Razaviyayn)
 Spandan Senapati (started Fall'23, co-advised with Haipeng Luo)
 Tianyi Zhou (started Fall'23, co-advised with Robin Jia)

Undergraduate and Masters students

Dutch Hansen
 You Qi Huang
 Anish Jayant
 Jung Whan Lee
 Natalie Abreu (graduated in Fall'23, now Ph.D. student at Harvard)
 Aditya Prasad (graduated in Fall'24, now Ph.D. student at the University of Chicago)
 Kameron Shahabi (graduated in Fall'24, now Ph.D. student at the University of Washington)
 Qilin Ye (graduated in Fall'24, now M.S. student at Duke)
 Devin Martin (SURE program intern in Summer'22, mentored by Bhavya Vasudeva)

TEACHING AT USC

Theory of Machine Learning: Fall'21, Fall'23
Machine Learning: Fall'22, Spring'24
Computational Perspectives on the Frontiers of Learning: Spring'23

PROFESSIONAL
ACTIVITIES

Reviewing for Grants

National Science Foundation (NSF) Panelist, 2022, 2023
Swiss National Science Foundation Reviewer, 2024

Reviewing for Journals

IEEE BITS “Special Issue on Generative Models” 2025 (co-editor)
IEEE Transactions on Information Theory, 2023
Journal of Machine Learning Research (JMLR), 2021, 2022
SIAM Journal of Computing, 2020, 2021

Conference Program Committee/Reviewing

Foundations of Responsible Computing (FORC), 2025 (PC member)
Conference on Learning Theory (COLT), 2019-2024, 2025 (area chair)
Algorithmic Learning Theory (ALT), 2022-2023, 2024-2025 (area chair)
International Conference on Learning Representations (ICLR), 2021-2023, 2025 (area chair)
International Conference on Machine Learning (ICML), 2019-2023, 2025 (area chair)
Neural Information Processing Systems (NeurIPS), 2019-2021, 2023, 2025 (area chair)
International Conference on Artificial Intelligence and Statistics (AISTATS), 2021
AAAI Conference on Artificial Intelligence (AAAI), 2020
Symposium on Theory of Computing (STOC), 2020-2022, 2024, 2025
Foundations of Computer Science (FOCS), 2020, 2023, 2024, 2025 (PC member)
Symposium on Discrete Algorithms (SODA), 2019, 2022, 2024, 2025
Innovations in Theoretical Computer Science (ITCS), 2020, 2023

OUTREACH

Learning Theory Alliance (LeT-All), Workshop Committee

- Organized day-long virtual mentoring workshop on “Day-to-day life of an ML/theory researcher” on June 4 2024, attended by around 300 students.
- Organized day-long virtual mentoring workshop on “Communicating your research” on October 26 2023, attended by around 200 students.

High-school students hosted in the summer (in collaboration with USC Viterbi K-12 Center)

Janna Audrey Doratan (Summer’24, mentored by Devansh Gupta)
Angela Zhuang (Summer’23, mentored by Julian Asilis)
Jayron Martinez (Summer’22, mentored by Siddhartha Devic)
Luke Pratt (Summer’22, mentored by Bhavya Vasudeva)

Talks for high-school students

Fair & Robust Artificial Intelligence

- As part of USC SHINE program for high-schoolers, June 2022
- Los Angeles County Office of Education (LACOE) CS speaker series, February 2023

INVITED TALKS
& SEMINARS

Using Algorithms to Understand Transformers (and Using Transformers to Understand Algorithms)

- Workshop on Theoretical Perspectives on LLMs, UCSD, March 2025
- Simons Institute Workshop on Transformers as a Computational Model, September 2024
- Google Research Mountain View, November 2024
- University of British Columbia CAIDA Seminar, December 2024

Discovering Data Structures: Nearest Neighbor Search and Beyond

- Information Theory and Applications, San Diego, February 2025

Transformers Learn Higher-Order Optimization Methods for In-Context Learning

- USC Symposium on Frontiers of Generative AI in Science and Society, March 2024
- Information Theory and Applications, San Diego, February 2024

Memory as a Lens to Understand Efficient Learning and Optimization

- UC Berkeley Theory Seminar, October 2024
- Workshop on Computational/Statistical Gaps in Learning & Optimization, UCLA, February 2024
- Indian Institute of Technology, Delhi, December 2023
- University of Maryland, College Park, September 2023
- Microsoft Research, NYC, September 2023

A multigroup perspective to go beyond loss minimization in ML

- BIRS Workshop on New Directions in Machine Learning Theory, Banff, October 2024

From Anomaly Detection to Robust ML: Multicalibration as an Algorithmic Paradigm

- Amazon Research, Palo Alto, July 2023

Fairness in Matching under Uncertainty

- Information Theory and Applications, San Diego, February 2023

Sample Amplification: Increasing Dataset Size even when Learning is Impossible

- IISA Annual Meeting, December 2024

- Simons Institute Seminar, September 2024

- USC Probability Seminar, October 2023

- Learning Theory Alliance Mentoring Workshop, February 2021

- Simons Institute, Learning in High Dimensions Program, September 2020

- Neurips Machine Learning with Guarantees Workshop, December 2019

Memory-sample Tradeoffs for Continuous Optimization and Learning

- MIT Theory Lunch, October 2020

- NYU Theory Seminar, November 2019

- EPFL Theory Seminar, November 2019

- Northwestern Junior Theorists Workshop, November 2019

- University of Washington Theory Lunch, October 2019

- Cornell ORIE Young Researchers Workshop, October 2019

- Google Research, Mountain View, August 2019

Prediction with a Short Memory

- ETH Zurich Institute for Theoretical Studies, November 2019

- China Theory Week, Tsinghua University, September 2018

- Google Mountain View Algorithms TechTalk, March 2018

Orthogonalized ALS: Theoretically Principled Tensor Factorization for Practical Use

- SIAM Annual Meeting, Portland, July 2018