

# CSCI 567: Machine Learning

Vatsal Sharan

Fall 2022

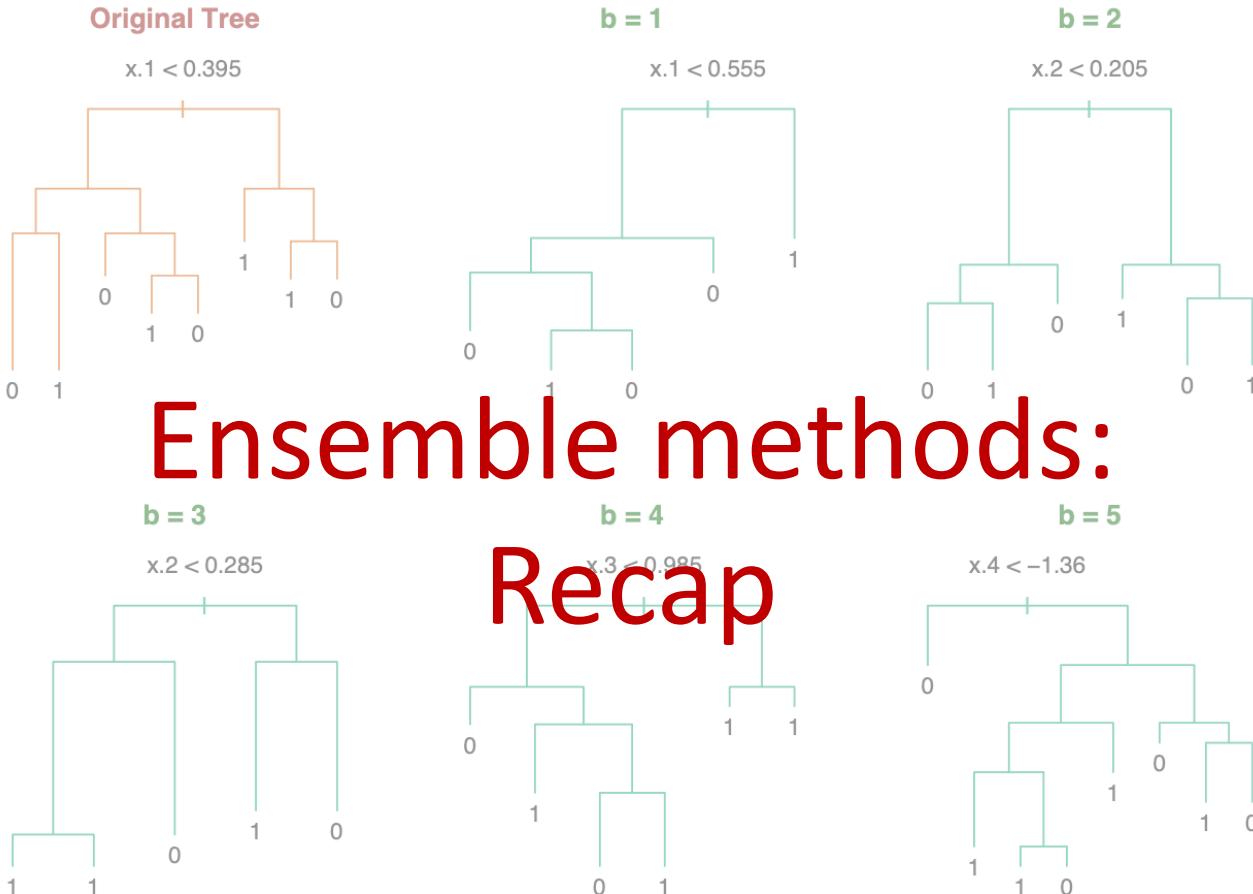
Lecture 9, Nov 3



**USC** University of  
Southern California

# Administrivia

- Project details are out
  - Make groups (of 4) by Nov 11, minimum team size is 3.
  - Q5 on HW4 will help you get started on it.
  - We'll give an overview of the project and general tips in today's discussion.
- HW4 is due in about two weeks (Nov 16 at 2pm).
  - We'll release another question on PCA tomorrow.
- Today's plan:
  - Finish ensemble methods
  - Unsupervised learning:
    - PCA
    - Clustering



# Ensemble methods

- Bagging
- Random forests
- Boosting: Basics
- Adaboost
- Gradient boosting

# Bagging

Collect  $T$  subsets each of some fixed size (say  $m$ ) by sampling with replacement from training data.

Let  $f_t(\mathbf{x})$  be the classifier (such as a decision tree) obtained by training on the subset  $t \in \{1, \dots, T\}$ . Then the aggregated classifier  $f_{agg}(\mathbf{x})$  is given by:

$$f_{agg}(\mathbf{x}) = \begin{cases} \frac{1}{T} \sum_{t=1}^T f_t(\mathbf{x}) & \text{for regression,} \\ \text{sign}\left(\frac{1}{T} \sum_{t=1}^T f_t(\mathbf{x})\right) = \text{Majority Vote}\{f_t(\mathbf{x})\}_{t=1}^T & \text{for classification.} \end{cases}$$

- Reduces overfitting (i.e., variance)
- Can work with any type of classifier (here focus on trees)
- Easy to parallelize (can train multiple trees in parallel)
- But loses on interpretability to single decision tree (true for all ensemble methods..)

# Ensemble methods

- Bagging
- Random forests
- Boosting: Basics
- Adaboost
- Gradient boosting

## Random forests

Random forests: When growing a tree on a bootstrapped (i.e. subsampled) dataset, before each split select  $k \leq d$  of the  $d$  input variables at random as candidates for splitting.

When  $k = d \rightarrow$  same as Bagging

When  $k < d \rightarrow$  Random forests

$k$  is a hyperparameter, tuned via cross-validation

# Ensemble methods

- Bagging
- Random forests
- **Boosting: Basics**
- Adaboost
- Gradient boosting

## Boosting: Idea

The boosted predictor is of the form  $f_{boost}(\mathbf{x}) = \text{sign}(h(\mathbf{x}))$ , where,

$$h(\mathbf{x}) = \sum_{t=1}^T \beta_t f_t(\mathbf{x}) \text{ for } \beta_t \geq 0 \text{ and } f_t \in \mathcal{F}.$$

The goal is to minimize  $\ell(h(\mathbf{x}), y)$  for some loss function  $\ell$ .

Q: We know how to find the best predictor in  $\mathcal{F}$  on some data, but how do we find the best weighted combination  $h(\mathbf{x})$ ?

A: Minimize the loss by a *greedy approach*, i.e. find  $\beta_t, f_t(\mathbf{x})$  one by one for  $t = 1, \dots, T$ .

Specifically, let  $h_t(\mathbf{x}) = \sum_{\tau=1}^t \beta_\tau f_\tau(\mathbf{x})$ . Suppose we have found  $h_{t-1}(\mathbf{x})$ , how do we find  $\beta_t, f_t(\mathbf{x})$ ?

Find the  $\beta_t, f_t(\mathbf{x})$  which minimizes the loss  $\ell(h_t(\mathbf{x}), y)$ .

Different loss function  $\ell$  give different boosting algorithms.

$$\ell(h(\mathbf{x}), y) = \begin{cases} (h(\mathbf{x}) - y)^2 & \rightarrow \text{Least squares boosting,} \\ \exp(-h(\mathbf{x})y) & \rightarrow \text{AdaBoost.} \end{cases}$$

# Ensemble methods

- Bagging
- Random forests
- Boosting: Basics
- Adaboost
- Gradient boosting

# AdaBoost: Full algorithm

Given a training set  $S$  and a base algorithm  $\mathcal{A}$ , initialize  $D_1$  to be uniform

For  $t = 1, \dots, T$

- obtain a weak classifier  $f_t(\mathbf{x}) \leftarrow \mathcal{A}(S, D_t)$
- calculate the weight  $\beta_t$  of  $f_t(\mathbf{x})$  as

$$\beta_t = \frac{1}{2} \ln \left( \frac{1 - \epsilon_t}{\epsilon_t} \right) \quad (\beta_t > 0 \Leftrightarrow \epsilon_t < 0.5)$$

where  $\epsilon_t = \sum_{i:f_t(\mathbf{x}_i) \neq y_i} D_t(i)$  is the weighted error of  $f_t(\mathbf{x})$ .

- update distributions

$$D_{t+1}(i) \propto D_t(i) e^{-\beta_t y_i f_t(\mathbf{x}_i)} = \begin{cases} D_t(i) e^{-\beta_t} & \text{if } f_t(\mathbf{x}_i) = y_i \\ D_t(i) e^{\beta_t} & \text{else} \end{cases}$$

Output the final classifier  $f_{boost} = \operatorname{sgn} \left( \sum_{t=1}^T \beta_t f_t(\mathbf{x}) \right)$

# Adaboost: Example

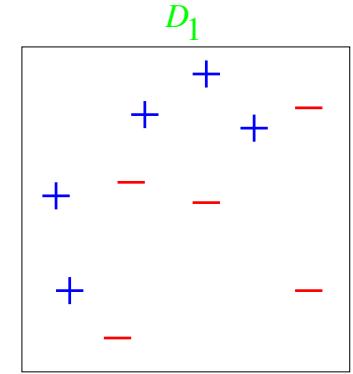
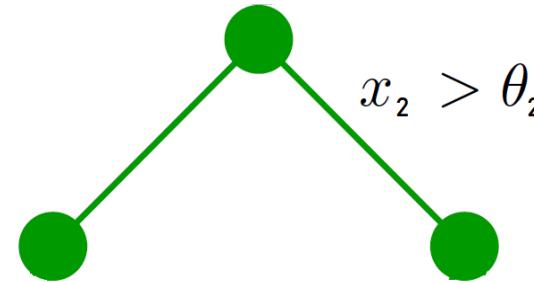
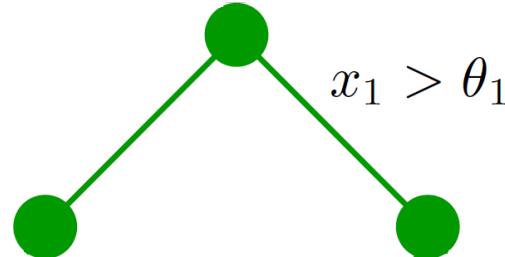
Put more weight on difficult to classify instances and less on those already handled well

New weak learners are added sequentially that focus their training on the more difficult patterns

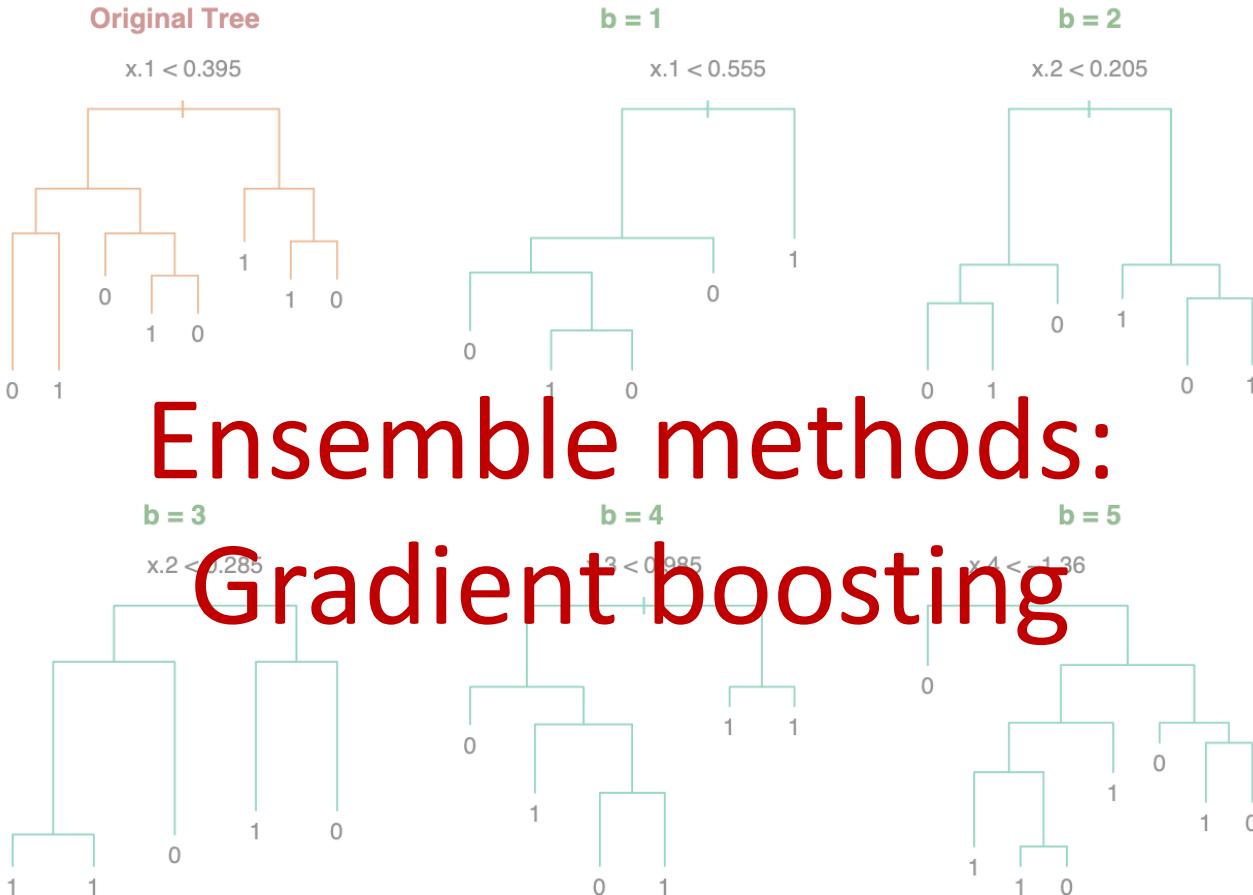
10 data points in  $\mathbb{R}^2$

The size of + or - indicates the weight, which starts from uniform  $D_1$

Base algorithm is decision stump:



Go through the calculations in the example to make sure you understand the algorithm



**Ensemble methods:  
Gradient boosting**

# Ensemble methods

- Bagging
- Random forests
- Boosting: Basics
- Adaboost
- Gradient boosting

# Gradient Boosting

Recall  $h_t(\mathbf{x}) = \sum_{\tau=1}^t \beta_\tau f_\tau(\mathbf{x})$ . For Adaboost (exponential loss), given  $h_{t-1}(\mathbf{x})$ , we found what  $f_t(\mathbf{x})$  should be.

Gradient boosting provides an iterative approach for general (any) loss function  $\ell(h(\mathbf{x}), y)$ :

- For all training datapoints  $(\mathbf{x}_i, y_i)$  find the gradient

$$r_i = \left[ \frac{\delta \ell(h(\mathbf{x}_i), y_i)}{\delta h(\mathbf{x}_i)} \right]_{h(\mathbf{x}_i)=h_{t-1}(\mathbf{x}_i)}$$

how should predictions change "locally" to reduce loss?

- Use the weak learner to find  $f_t$  which fits  $(\mathbf{x}_i, r_i)$  as well as possible:

$$f_t = \operatorname{argmin}_{f \in \mathcal{F}} \sum_{i=1}^n (r_i - f(\mathbf{x}_i))^2.$$

fit a model to  $r_i$

- Update  $h_t(\mathbf{x}) = h_{t-1}(\mathbf{x}) + \eta f_t(\mathbf{x})$ , for some step size  $\eta$ .

add model which improves loss locally

# Gradient Boosting

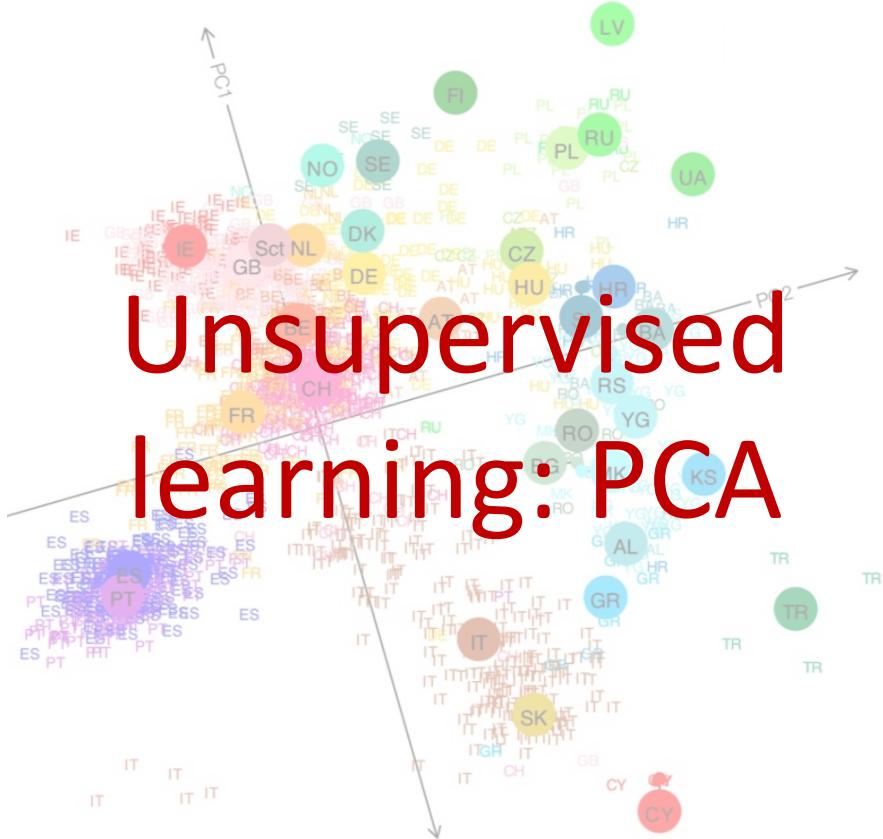
Usually we add some regularization term to prevent overfitting (penalize the size of the tree etc.)

Gradient boosting is extremely successful!!

A variant **XGBoost** is one of the most popular algorithms for **structured data** (tables etc. with numbers and categories where each feature typically has some meaning, unlike images or text).

(for e.g. during Kaggle competitions back in 2015, 17 out of 29 winning solutions used XGBoost)

# Unsupervised learning: PCA



# A simplistic taxonomy of ML

## **Supervised learning:**

Aim to predict outputs of future datapoints

## **Unsupervised learning:**

Aim to discover hidden patterns and explore data

## **Reinforcement learning:**

Aim to make sequential decisions

# Principal Component Analysis (PCA)

- Introduction
- Formalizing the problem
- How to use PCA, and examples
- Solving the PCA optimization problem
- Conclusion

# Acknowledgement & further reading

Our presentation is closely based on Gregory Valiant's notes for CS168 at Stanford.

<https://web.stanford.edu/class/cs168/l/l7.pdf>

<https://web.stanford.edu/class/cs168/l/l8.pdf>

You can refer to these notes for further reading.

Also review our Linear algebra Colab notebooks:

[Part 1](#)

[Part 2](#)

# Dimensionality reduction & PCA

We'll start with a simple and fundamental unsupervised learning problem: **dimensionality reduction**.

**Goal:** reduce the dimensionality of a dataset so that

- it is easier to visualize and discover patterns
- it takes less time and space to process for any downstream application (classification, regression, etc)
- noise is reduced
- ...

There are many approaches, we focus on a linear method: **Principal Component Analysis (PCA)**.

# PCA: Motivation

Consider the following dataset:

- 17 features, each represents the **average consumption of some food**
- 4 data points, each represents some **country**.

*What can you tell?*

Hard to say anything looking at all these 17 features.

Alcoholic drinks  
Beverages  
Carcase meat  
Cereals  
Cheese  
Confectionery  
Fats and oils  
Fish  
Fresh fruit  
Fresh potatoes  
Fresh Veg  
Other meat  
Other Veg  
Processed potatoes  
Processed Veg  
Soft drinks  
Sugars

	England	N Ireland	Scotland	Wales
Alcoholic drinks	375	135	458	475
Beverages	57	47	53	73
Carcase meat	245	267	242	227
Cereals	1472	1494	1462	1582
Cheese	105	66	103	103
Confectionery	54	41	62	64
Fats and oils	193	209	184	235
Fish	147	93	122	160
Fresh fruit	1102	674	957	1137
Fresh potatoes	720	1033	566	874
Fresh Veg	253	143	171	265
Other meat	685	586	750	803
Other Veg	488	355	418	570
Processed potatoes	198	187	220	203
Processed Veg	360	334	337	365
Soft drinks	1374	1506	1572	1256
Sugars	156	139	147	175

Picture from [here](#)  
See [this](#) for more details

# PCA: Motivation

PCA can help us! The **projection of the data onto its first principal component**:

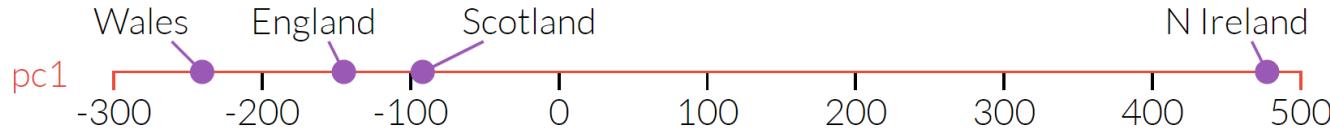


i.e. we reduce the dimensionality from 17 to just 1.

Now one data point is clearly different from the rest!

# PCA: Motivation

PCA can help us! The **projection of the data onto its first principal component (PC1)**:



i.e. we reduce the dimensionality from 17 to just 1.

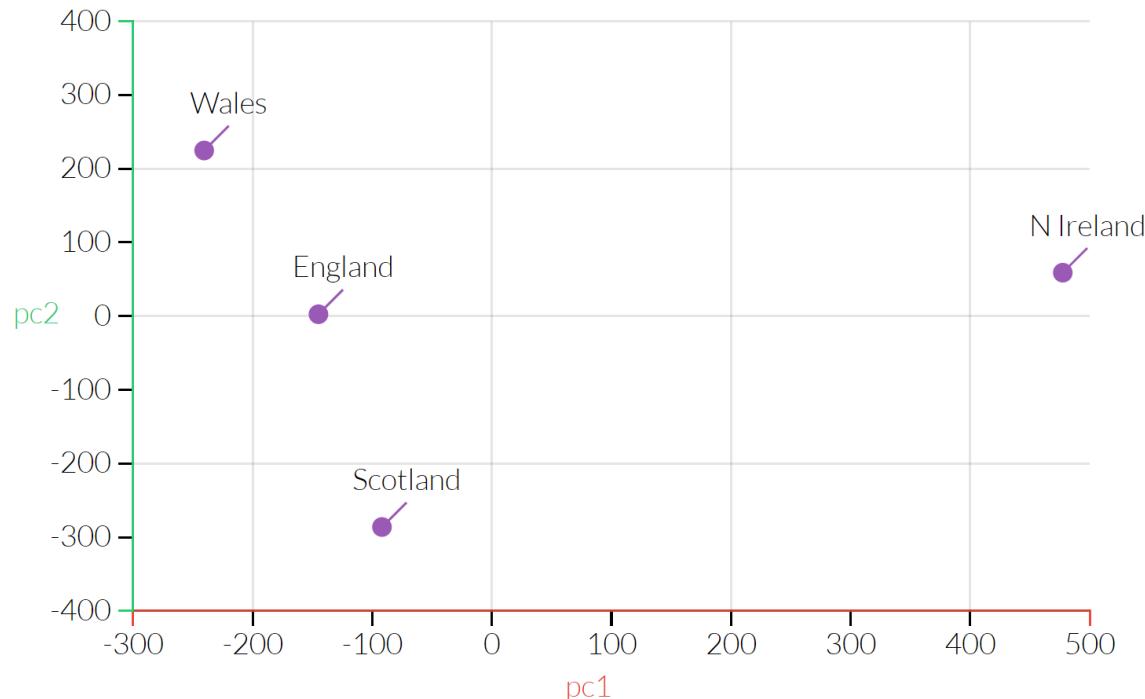
Now one data point is clearly different from the rest!

That turns out to be data from Northern Ireland,  
the only country not on the island of Great Britain out of the 4 samples.

Can also interpret components: PC1 tells us that the Northern Irish eat more grams of fresh potatoes and fewer of fresh fruits and alcoholic drinks.

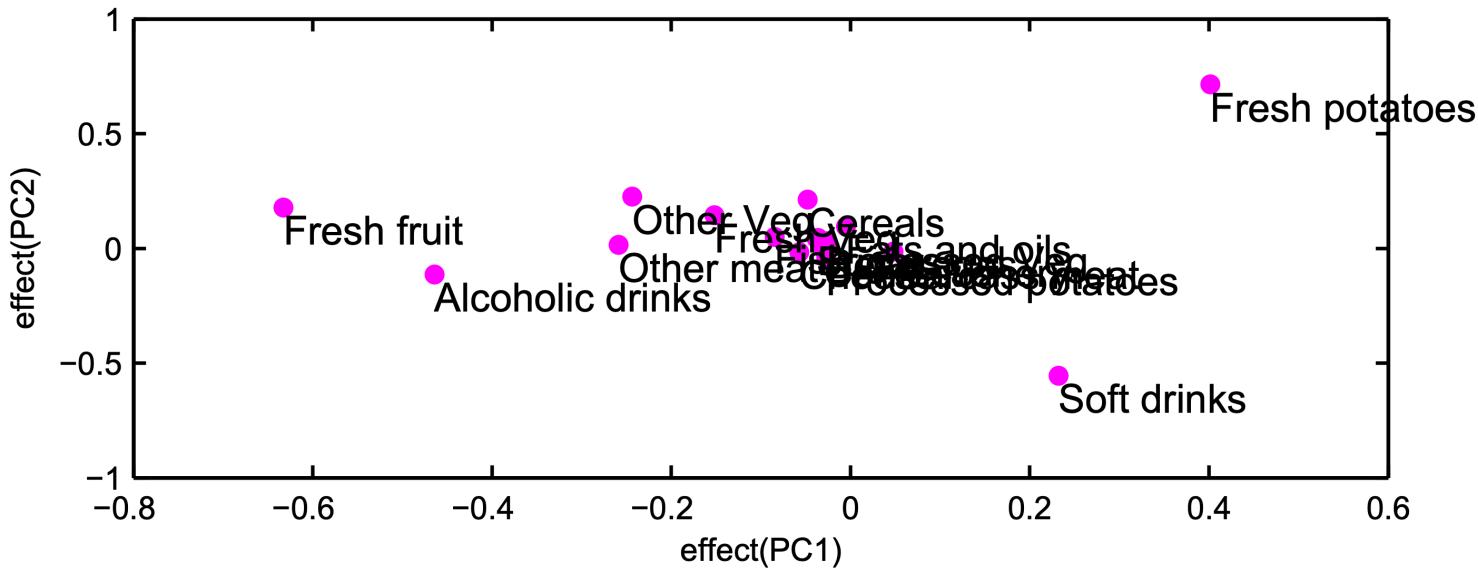
# PCA: Motivation

We can find the **second (and more)** principal components of the data too:



# PCA: Motivation

And the components themselves are interpretable too:



See [this](#) for more details

# Principal Component Analysis (PCA)

- Introduction
- Formalizing the problem
- How to use PCA, and examples
- Solving the PCA optimization problem
- Conclusion

## High-level goal

prev. e.g.

$n=4$

$d=17$

Suppose we have a dataset of  $n$  datapoints  $x_1, x_2, \dots, x_n \in \mathbb{R}^d$ .

The high level goal of PCA is to find a set of  $k$  principal components (PCs) /principal vectors  $v_1, \dots, v_k \in \mathbb{R}^d$  such that for each  $x_i$ ,

$$x_i \approx \sum_{j=1}^k \alpha_{ij} v_j$$

"principal food consumption vectors".  
↳ food consumption for some country

for some coefficients  $\alpha_{ij} \in \mathbb{R}$ .

Explain the data as different linear combinations of some PCs

# Preprocessing the data

- Before we apply PCA, we usually preprocess the data to center it

$$\text{Let } \bar{x} = \frac{1}{n} \sum_{i=1}^n x_i$$

$$\text{Then set } \tilde{x}_i = x_i - \bar{x}$$

Assume data is centered ( $\sum x_i = 0$ )

- In many applications, it is also important to scale each coordinate properly. This is especially true if the coordinates are in different units or scales.

for all  $j \in [d]$ , divide  $j$ th coordinate of each point by  $\sqrt{\sum_{i=1}^n x_{ij}^2}$

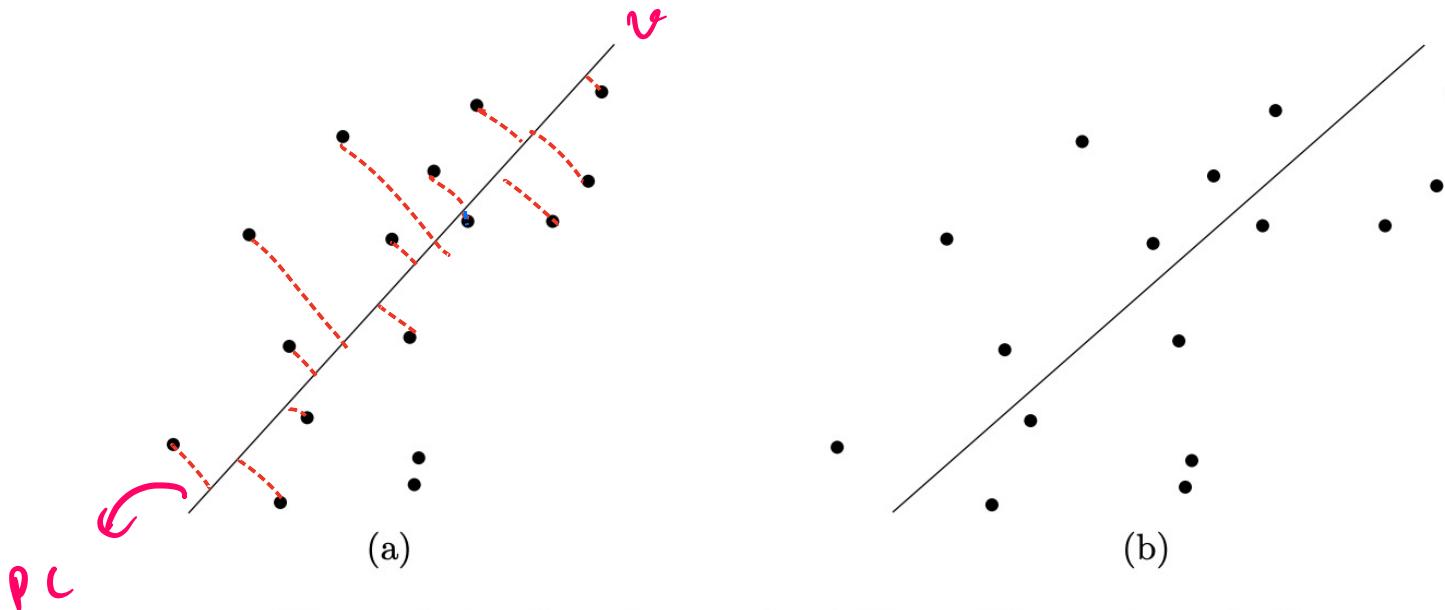


Figure 3: Scaling the  $x$ -axis yields a different best-fit line.

# Objective function for PCA

Key difference from supervised learning problems:

No labels given, which means no ground-truth to measure the quality of the answer!

However, we can still write an optimization problem based on our high-level goal.

For clarity, we first discuss the special case of  $k = 1$ .

Optimization problem for finding the 1<sup>st</sup> principal component  $v_1$ :

$$v_1 = \arg \min_{\mathbf{v}: \|v\|_2=1} \sum_{i=1}^n ((\text{distance between } x_i \text{ & line spanned by } v)^2)$$

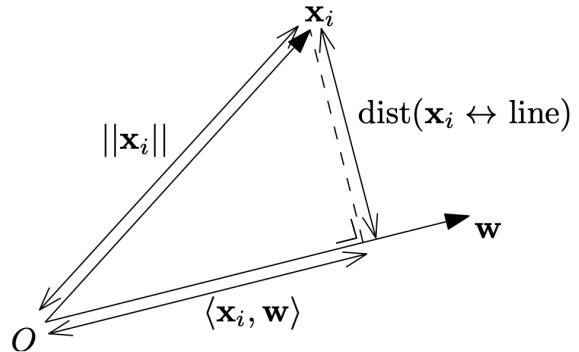


Figure 4: The geometry of the inner product with a unit length vector,  $\mathbf{w}$ .

$$(\text{dist}(\mathbf{x}_i \leftrightarrow \text{line spanned by } \mathbf{v}))^2 + \langle \mathbf{x}_i, \mathbf{v} \rangle^2 = \|\mathbf{x}_i\|_2^2$$

$\|\mathbf{x}_i\|_2^2$  is a constant, independent of choice of  $\mathbf{v}$ .

∴ original objective is equivalent to

$$v_1 = \arg \min_{v} \sum_{i=1}^n$$

$$v : \|v\|_2 = 1$$

$$\sum_{i=1}^n \langle x_i, v \rangle^2$$

An example:

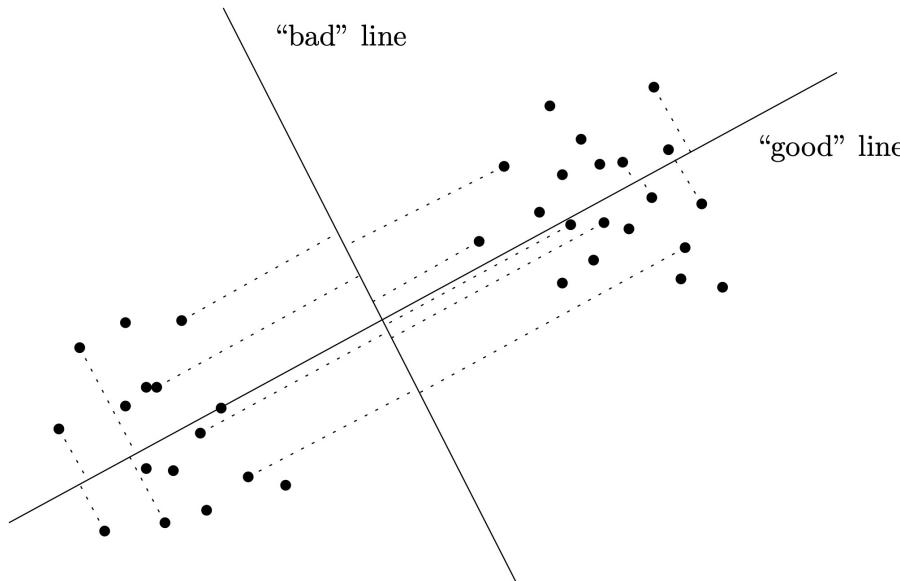


Figure 5: For the good line, the projection of the points onto the line keeps the two clusters separated, while the projection onto the bad line merges the two clusters.

# Objective function for larger values of $k$

The generalization of the original formulation for general  $k$  is to find a  $k$ -dimensional subspace  $S$  such that the points are as close to it as possible:

$$S = \underset{k\text{-dim subspaces } S}{\operatorname{argmin}} \sum_{i=1}^n (\text{distance between } \mathbf{x}_i \text{ and subspace } S)^2$$

By the same reasoning as for  $k = 1$ , this is equivalent to,

$$S = \underset{k\text{-dim subspaces } S}{\operatorname{argmax}} \sum_{i=1}^n (\text{length of } \mathbf{x}_i \text{'s projection on } S)^2 \quad - \textcircled{1}$$

It is useful to think of the subspace  $S$  as the *span* of  $k$  *orthonormal vectors*  $\mathbf{v}_1, \dots, \mathbf{v}_k \in \mathbb{R}^d$ .

Recall, vectors  $\mathbf{v}_1, \dots, \mathbf{v}_k$  are orthonormal if

$$\textcircled{1} \quad \|\mathbf{v}_i\|_2 = 1 \quad \forall i \in [k]$$

$$\textcircled{2} \quad \langle \mathbf{v}_i, \mathbf{v}_j \rangle = 0 \quad \forall i \neq j$$

| C-g.  
standard basis vectors  
 $\mathbf{e}_1 = (1, 0, 0, \dots, 0)$   
 $\mathbf{e}_2 = (0, 1, 0, \dots, 0)$   
⋮

Def. Span of a collection  $v_1, \dots, v_k \in \mathbb{R}^d$

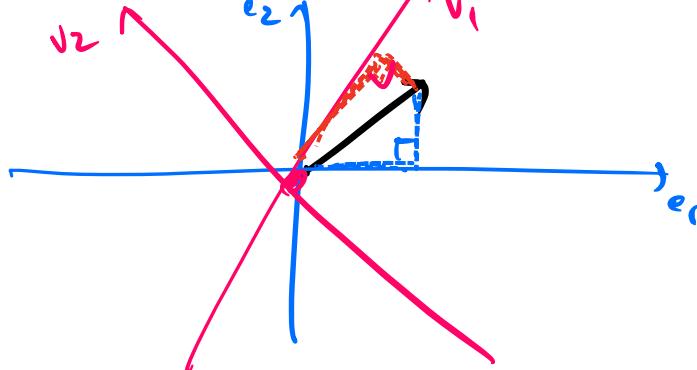
is all their linear combinations  $\left\{ \sum_{j=1}^k d_j v_j : d_1, \dots, d_k \in \mathbb{R} \right\}$

Example,

- $k = 1$ , span is line through the origin.
- $k = 2$ , if  $v_1, v_2$  are linearly independent, the span is a plane through the origin, and so on.

fact about orthonormal vectors,

$$(\text{length of } x_i \text{'s projection on } \text{span}(v_1, \dots, v_k))^2 = \sum_{j=1}^k \langle x_i, v_j \rangle^2 - \textcircled{2}$$



Combining ① & ②

Formal problem solved by PCA:

Given  $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^d$  and a parameter  $k \geq 1$ , compute orthonormal vectors  $\mathbf{v}_1, \dots, \mathbf{v}_k \in \mathbb{R}^d$  to maximize,

$$\sum_{i=1}^n \sum_{j=1}^k \langle \mathbf{x}_i, \mathbf{v}_j \rangle^2.$$

Equivalent view:

- Pick  $\mathbf{v}_1$  to be the variance maximizing direction.
- Pick  $\mathbf{v}_2$  to be the variance maximizing direction, orthogonal to  $\mathbf{v}_1$ .
- Pick  $\mathbf{v}_3$  to be the variance maximizing direction, orthogonal to  $\mathbf{v}_1$  and  $\mathbf{v}_2$ , and so on.

# Principal Component Analysis (PCA)

- Introduction
- Formalizing the problem
- How to use PCA, and examples
- Solving the PCA optimization problem
- Conclusion

# Using PCA for data compression and visualization

**Input:**  $n$  datapoints  $x_1, x_2, \dots, x_n \in \mathbb{R}^d$ , #components  $k$  we want

**Step 1** Perform PCA to get top  $k$  principal components  $v_1, \dots, v_k \in \mathbb{R}^d$ .

**Step 2** For each datapoint  $x_i$ , define its “ $v_1$ -coordinate” as  $\langle x_i, v_1 \rangle$ , its “ $v_2$ -coordinate” as  $\langle x_i, v_2 \rangle$ . Therefore we associate  $k$  coordinates to each datapoint  $x_i$ , where the  $j$ -th coordinate denotes the extent to which  $x_i$  points in the direction of  $v_j$ .

**Step 3** We now have a new “compressed” dataset where each datapoint is  $k$ -dimensional. For visualization, we can plot the point  $x_i$  in  $\mathbb{R}^k$  as the point  $(\langle x_i, v_1 \rangle, \langle x_i, v_2 \rangle, \dots, \langle x_i, v_k \rangle)$ .

Going back to high-level goal,

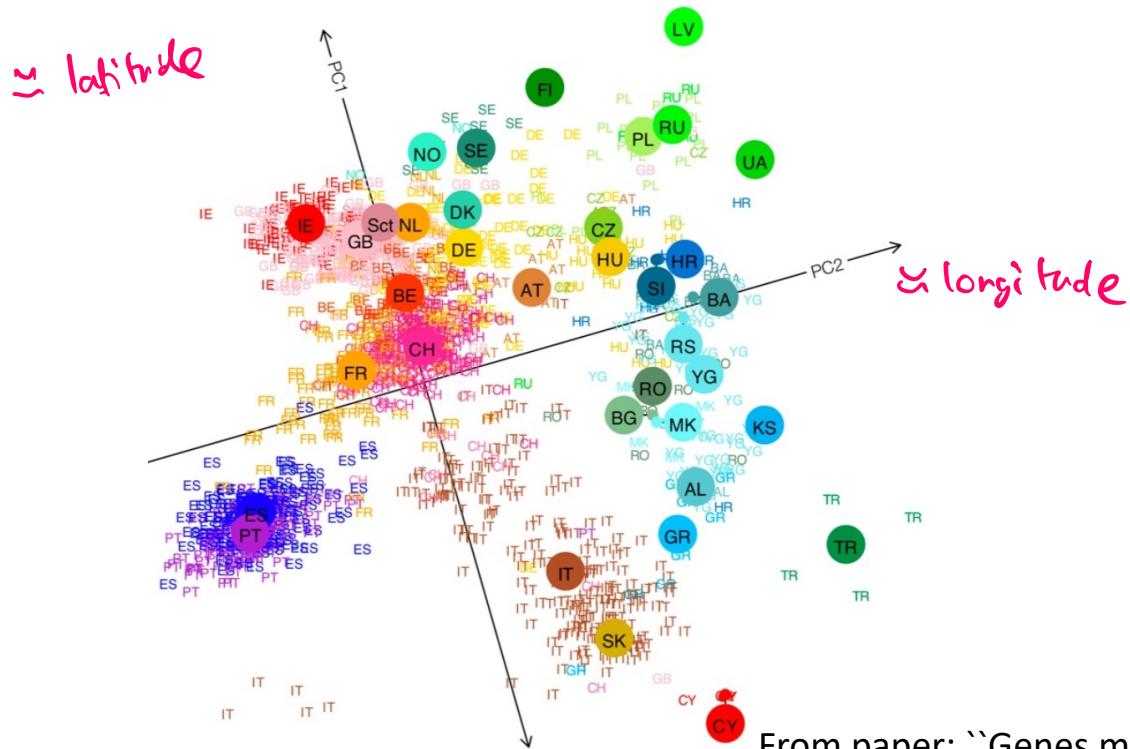
$$x_i \approx \sum_{j=1}^k \langle x_i, v_j \rangle v_j.$$

## Visualization example: Do Genomes Encode Geography?

Dataset: genomes of 1,387 Europeans (each individual's genotype at 200,000 locations in the genome)  
 $n = 1387, d \approx 200,000$

## Project the datapoints onto top 2 PCs

Plot shown below; looks remarkably like the map of Europe!



From paper: "Genes mirror geography within Europe" Novembre et al., Nature'08

# Compression example: Eigenfaces

Dataset: 256x256 ( $\approx$ 65K pixels) dimensional images of about 2500 faces, all framed similarly

$$n = 2500, d \approx 65,000$$

We can represent each image with high accuracy using only 100-150 principal components!

The principal components (called *eigenfaces* here) are themselves interpretable too!



**Figure 2.** Seven of the eigenfaces calculated from the input images of Figure 1.

# Principal Component Analysis (PCA)

- Introduction
- Formalizing the problem
- How to use PCA, and examples
- Solving the PCA optimization problem
- Conclusion

# How to solve the PCA optimization problem?

Consider  $k=1$ ,

$$V_1 = \underset{v}{\operatorname{arg\,max}} \sum_{i=1}^n \langle f_i, v \rangle^2$$

$v: \|v\|_2 = 1$

$$X = \begin{bmatrix} f_1^\top \\ f_2^\top \\ \vdots \\ f_n^\top \end{bmatrix} \in \mathbb{R}^{n \times d}$$

$$\therefore \text{for any } v \in \mathbb{R}^d, \quad f^T v = \begin{bmatrix} f_1^\top v \\ f_2^\top v \\ \vdots \\ f_n^\top v \end{bmatrix} \in \mathbb{R}^n$$

$$\begin{aligned}\therefore \sum_{i=1}^n \langle x_i, v \rangle^2 &= \|xv\|_2^2 \\ &= (xv)^T (xv) \\ &= v^T \underbrace{x^T x}_{d+n \times d} v\end{aligned}$$

$\therefore$  for  $A = \underbrace{x^T x}_{d+n \times d} \in \mathbb{R}^{d+d}$

$$v_1 = \underset{\|v\|_2 = 1}{\text{argmax}} \quad v^T A v$$

$x^T x$ : covariance matrix of data (assuming data is centered)

$$A = \begin{bmatrix} 1 & 1 & \dots & 1 \\ f_1 & f_2 & \dots & f_n \\ 1 & 1 & \dots & 1 \end{bmatrix} \begin{bmatrix} 1 & f_1^2 & \dots \\ - & f_2^2 & \dots \\ \vdots & \vdots & \ddots \\ - & f_n^2 & \dots \end{bmatrix}$$

$$\begin{aligned}A_{11} &= \sum_{i=1}^n x_{i,1}^2 \rightarrow \text{variance of 1st coordinate} \\ A_{12} &= \sum_{i=1}^n x_{i,1} f_{i,2} \rightarrow \text{covariance b/w 1st \& 2nd coordinates}\end{aligned}$$

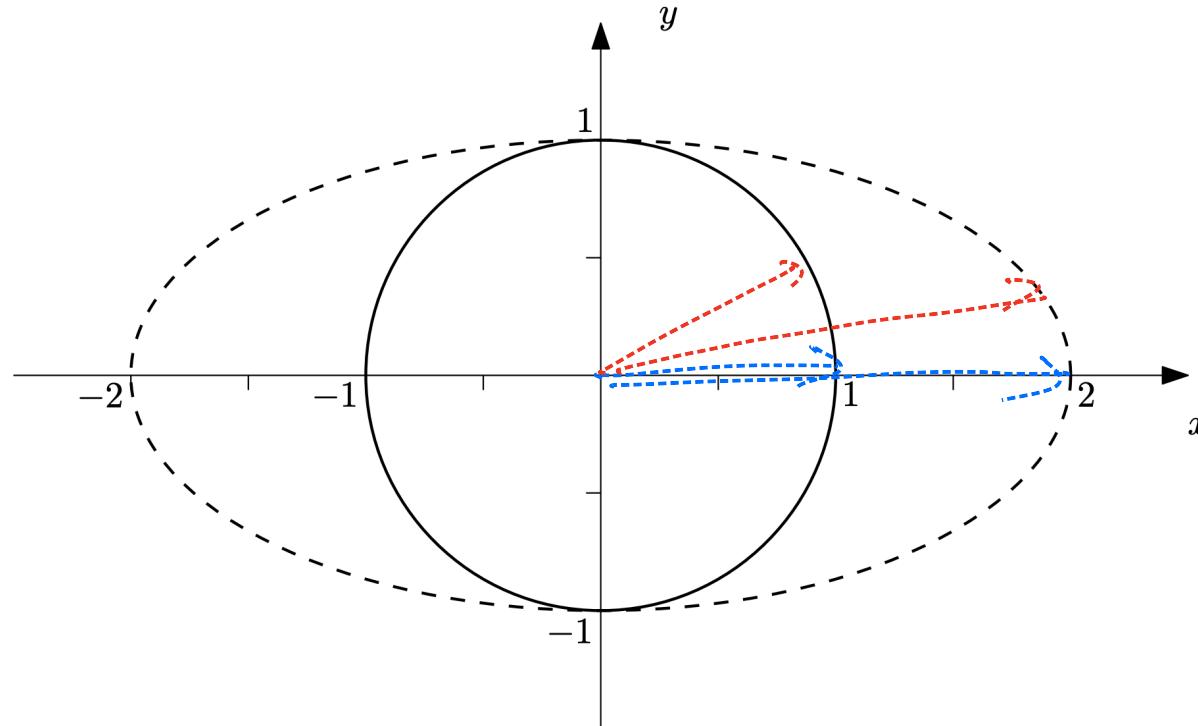
## The diagonal case

Let's solve  $\operatorname{argmax}_{v: \|v\|_2=1} v^T A v$  for the special case where  $A$  is a diagonal matrix.

$$A = \begin{pmatrix} d_1 & & & \\ & d_2 & \dots & 0 \\ & & \ddots & \\ 0 & & & d_d \end{pmatrix} \quad \text{where } d_1 \geq d_2 \geq \dots \geq d_d \geq 0$$

Any  $d \times d$  matrix  $A$  can be thought of as a function that maps points in  $\mathbb{R}^d$  back to points in  $\mathbb{R}^d$ :  
 $v \mapsto Av$ .

The matrix  $\begin{pmatrix} 2 & 0 \\ 0 & 1 \end{pmatrix}$  maps  $(x, y)$  to  $(2x, y)$ :



Points on circle  $\{(x, y) : x^2 + y^2 = 1\}$  are mapped to the ellipse  $\{(x, y) : \left(\frac{x}{2}\right)^2 + y^2 = 1\}$ .

So what direction  $v$  should maximize  $v^T A v$  for diagonal  $A$ ?

It should be the direction of maximum stretch:

$v = e_1$  (where  $e_i = (1, 0, \dots, 0)$  is 1st standard basis vector)

(since  $d_1 \geq d_2 \dots \geq d_d$ )

Proof:

$$v^T A v = v^T (A v) = (v_1 \dots v_d) \cdot \begin{pmatrix} d_1 v_1 \\ d_2 v_2 \\ \vdots \\ d_d v_d \end{pmatrix} = \sum_{i=1}^d v_i^2 \cdot d_i$$

Since  $v$  is a unit vector,  $\sum_{i=1}^d v_i^2 = 1$

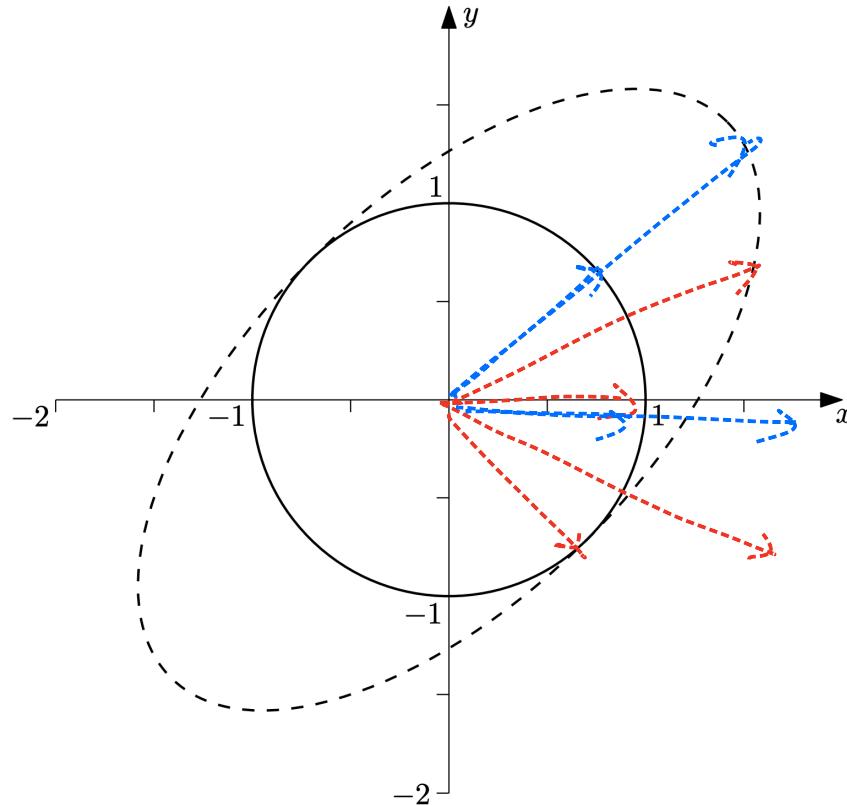
$\therefore$  Since  $d_1$  is largest, to max. set  $v_1 = 1$   
 $v_i = 0 \quad \forall i > 0.$   $\Rightarrow v = e_1.$

# Diagonals in disguise

Consider

$$\begin{aligned} A &= \begin{pmatrix} \frac{3}{2} & \frac{1}{2} \\ \frac{1}{2} & \frac{3}{2} \end{pmatrix} \xrightarrow{Q} \xrightarrow{D} \xrightarrow{Q^T} \\ &= \left( \begin{pmatrix} \frac{1}{\sqrt{2}} & -\frac{1}{\sqrt{2}} \\ \frac{1}{\sqrt{2}} & \frac{1}{\sqrt{2}} \end{pmatrix} \cdot \begin{pmatrix} 2 & 0 \\ 0 & 1 \end{pmatrix} \cdot \begin{pmatrix} \frac{1}{\sqrt{2}} & \frac{1}{\sqrt{2}} \\ -\frac{1}{\sqrt{2}} & \frac{1}{\sqrt{2}} \end{pmatrix} \right) \\ &\quad \text{rotates back by } 45^\circ \qquad \text{stretch} \qquad \text{rotate clockwise by } 45^\circ. \end{aligned}$$

$A$  still does nothing other than stretch out different orthogonal axes, possibly with these axes being a “rotated version” of the original ones.



The previous figure, rotated 45 degrees.

How do we formalize the concept of a rotation in high dimensions as a matrix operation?

Answer: Orthogonal matrix (also called orthonormal matrix).

An orthogonal matrix is a matrix  $\mathbf{Q}$  s.t. for all columns

$$\mathbf{q}_1, \dots, \mathbf{q}_d,$$

$$\|\mathbf{q}_i\|_2^2 = 1 \quad \forall i$$

$$\mathbf{q}_i^T \mathbf{q}_j = 0 \quad \forall i \neq j$$

Key properties:

$$① \mathbf{Q}^T \mathbf{Q} = \mathbf{I} \quad (\mathbf{Q}^{-1} = \mathbf{Q}^T)$$

$$② \|\mathbf{Qv}\|_2^2 = \|v\|_2^2$$

③ If  $\mathbf{Q}$  is orthogonal,  $\mathbf{Q}^T$  is also orthogonal.

①

$$\begin{pmatrix} -\mathbf{q}_1^T \\ -\mathbf{q}_2^T \\ -\mathbf{q}_d^T \end{pmatrix} \begin{pmatrix} 1 & 1 & 1 \\ 0 & 1 & 0 \\ 1 & 1 & 1 \end{pmatrix} = \begin{pmatrix} 1 & 1 & 1 \\ 0 & 1 & 0 \\ 1 & 1 & 1 \end{pmatrix}$$

②

$$\begin{aligned} \|\mathbf{Qv}\|_2^2 &= (\mathbf{Qv})^T (\mathbf{Qv}) \\ &= v^T \mathbf{Q}^T \mathbf{Q} v \\ &= v^T v = \|v\|_2^2 \end{aligned}$$

③

$$\begin{aligned} \mathbf{Q}^T \mathbf{Q} &= \mathbf{I} \\ \Rightarrow \mathbf{Q}(\mathbf{Q}^T \mathbf{Q}) &= \mathbf{Q} \\ \Rightarrow (\mathbf{Q} \mathbf{Q}^T) \mathbf{Q} &= \mathbf{Q} \Rightarrow \mathbf{Q} \mathbf{Q}^T = \mathbf{I} \end{aligned}$$

Recall that we want to find  $\mathbf{v}_1 = \operatorname{argmax}_{\mathbf{v}: \|\mathbf{v}\|_2=1} \mathbf{v}^T \mathbf{A} \mathbf{v}$ .

Now consider  $\mathbf{A}$  that can be written as  $\mathbf{A} = \mathbf{Q} \mathbf{D} \mathbf{Q}^T$  for an orthogonal matrix  $\mathbf{Q}$  and diagonal matrix  $\mathbf{D}$  with diagonal entries  $\lambda_1 \geq \lambda_2 \geq \lambda_3 \geq \dots \lambda_d \geq 0$ .

$$\mathbf{A} = \begin{array}{c} \boxed{\begin{matrix} 1 & | & | & | \\ | & | & \ddots & | \\ 0_1 & 0_2 & \dots & 0_d \\ | & | & | & | \end{matrix}} & \boxed{\begin{matrix} \lambda_1 & & & \\ & \lambda_2 & \dots & 0 \\ & 0 & \ddots & \vdots \\ & & & \lambda_d \end{matrix}} & \boxed{\begin{matrix} -0_1^2 \\ -0_2^2 \\ \vdots \\ -0_d^2 \end{matrix}} \\ \mathbf{Q} & \mathbf{D} & \mathbf{Q}^T \end{array} \left( \begin{matrix} \mathbf{Q} \mathbf{e}_1 \end{matrix} \right)$$

$\mathbf{Q}$  will only rotate, so doesn't effect maximization  $\mathbf{v}$

What is the direction which gets stretched the maximum?

(Informal answer) The maximum possible stretch by  $\mathbf{D}$  is  $\lambda_1$ . The direction of maximum stretch under  $\mathbf{D}$  is  $\mathbf{e}_1$ . Therefore, direction of maximum stretch under  $\mathbf{D} \mathbf{Q}^T$  is  $\mathbf{v}$  s.t.  $\mathbf{Q}^T \mathbf{v} = \mathbf{e}_1 \implies \mathbf{v} = (\mathbf{Q}^T)^{-1} \mathbf{e}_1 = \mathbf{Q} \mathbf{e}_1$ .

Claim : for  $A = Q D Q^T$

$$\text{argmax } v^T A v = Q e_1.$$

$$v: \|v\|_2 = 1$$

Proof : for  $v_i = Q e_1$

$$v_i^T A v_i = (Q e_1)^T A (Q e_1) = e_1^T \underbrace{Q^T}_{Q^T Q = I} \underbrace{Q D Q^T}_{Q^T Q = I} Q e_1 \\ = e_1^T D e_1 \\ = d_{11}$$

$$\text{Now for } v^T A v = \left( \begin{array}{c|c} \text{---} & 1 \\ \hline v^T & \end{array} \right) \left( \begin{array}{c|c} Q & \\ \hline \end{array} \right) \left( \begin{array}{c|c} d_{11} & \dots & d_{1n} \\ \hline 0 & \dots & 0 \end{array} \right) \left( \begin{array}{c|c} Q^T & \\ \hline \end{array} \right) \left( \begin{array}{c|c} \text{---} & 1 \\ \hline v & \end{array} \right)$$

$Q$  &  $Q^T$  preserve length  $\therefore Q^T Q = I$  or unit vectors.

$$\therefore v^T A v \leq d_{11}$$

$\therefore v_i = Q e_1$  maximizes  $v^T A v$ .

•

# General covariance matrices

Consider any covariance  $A = X^T X$ .

Linear algebra fact:

Any symmetric matrix  $A$  can be written as  $A = Q D Q^T$  for orthogonal matrix  $Q$  & diagonal matrix  $D$ .

If  $A = X^T X$ ,  $A$  is symmetric &  $D$  always has non-negative entries. Why?

$$v^T A v = v^T X^T X v = \|Xv\|_2^2 \geq 0$$

If  $D_{ii} < 0$ , then  $v = Q e_i$   $e^T Q D Q^T v < 0$ .

When  $k = 1$ , the solution to  $\operatorname{argmax}_{v: \|v\|_2=1} v^T A v$  is the first column of  $Q$ , where  $A = X^T X = Q D Q^T$  with  $Q$  orthogonal and  $D$  diagonal with sorted entries.

## General values of $k$

What is the solution to the PCA objective for general values of  $k$ ?

$$\sum_{i=1}^n \sum_{j=1}^k \langle \mathbf{x}_i, \mathbf{v}_j \rangle^2$$

Solution: Pick the first  $k$  columns of  $\mathbf{Q}$ , where the covariance  $\mathbf{X}^\top \mathbf{X} = \mathbf{Q} \mathbf{D} \mathbf{Q}^\top$  with  $\mathbf{Q}$  orthogonal and  $\mathbf{D}$  diagonal with sorted entries.

Since  $\mathbf{Q}$  is orthogonal, the first  $k$  columns of  $\mathbf{Q}$  are orthogonal vectors. These are called the top  $k$  principal components (PCs).

# Eigenvalues & eigenvectors

How to compute the top  $k$  columns of  $Q$  in the decomposition  $X^T X = Q D Q^T$ ?

Solution: Eigenvalue decomposition!

Def: An eigenvector of matrix  $A$  is a vector  $v$  that is stretched in the same direction by  $A$ , i.e.

$$Av = \lambda v \quad \text{for some } \lambda \in \mathbb{R}.$$

$\lambda$  is the eigen value.

Eigenvectors: axes of stretch in geometric intuition

Eigenvalues: stretch factors

When we write  $A = X^T X$  as  $A = Q D Q^T$ ,

→ rows of  $Q^T$  (columns of  $Q$ ) are eigenvectors of  $A$ .

→ diagonal entries of  $D$  are corresponding eigenvalues.

Proof :

i<sup>th</sup> column of  $Q$  is given by  $Q e_i$

$$\left( \begin{array}{c|c} \begin{matrix} 1 & 1 & \dots & 1 \\ 0 & 0 & \dots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \dots & 0 \end{matrix} & \begin{matrix} 0 \\ \vdots \\ 0 \end{matrix} \end{array} \right)$$

$$A(Q e_i) = Q D Q^T Q e_i = Q D e_i = Q d_i e_i = d_i (Q e_i).$$

∴ i<sup>th</sup> column of  $Q$  is eigenvector of  $A$  with eigenvalue  $d_i$ .

PCA boils down to computing the  $k$  eigenvectors of the covariance matrix  $X^T X$  that have the largest eigenvalues.

# Principal Component Analysis (PCA)

- Introduction
- Formalizing the problem
- How to use PCA, and examples
- Solving the PCA optimization problem
- Conclusion

## How many PCs to use?

For visualization, we usually choose  $k$  to be small and just pick the first few principal components.

In other applications such as compression, it is a good idea to plot the eigenvalues and see. A lot of data is close to being low rank, so the eigenvalues may decay and become small.

We can also choose the threshold based on how much variance we want to capture. Suppose we want to capture 90% of the variance in the data. Then we can pick  $k$  such that i.e.

$$\frac{\sum_{j=1}^k \lambda_j}{\sum_{j=1}^d \lambda_j} \geq 90\%$$

where  $\lambda_1 \geq \dots \geq \lambda_d$  are sorted eigenvalues.

Note:  $\sum_{j=1}^d \lambda_j = \text{trace}(\mathbf{X}^T \mathbf{X})$ , so no need to actually find all eigenvalues.

sum of diagonal entries

# When and why does PCA fail?

1. Data is not properly scaled/normalized.
2. Non-orthogonal structure in data: PCs are forced to be orthogonal, and there may not be too many orthogonal components in the data which are all interpretable.
3. Non-linear structure in data.

